



Bharatiya Vidya Bhavan's
Sardar Patel Institute of Technology
Bhavan's Campus, Munshi Nagar, Andheri (West), Mumbai-400058-India
(Autonomous College Affiliated to University of Mumbai)

Subject: NLP

Branch: AIML

Name	Jash Vora
UID	2022600065
Batch	A1
Course	NLP
Exp no.	8
Name of the Experiment	Retrieval Augmented Generation (RAG) based Gen-AI Tool
AIM	Build a Gen-AI tool using a vector database. Students will select a topic from the provided list (image).
Theory	<p>Retrieval Augmented Generation (RAG):</p> <p>Combines information retrieval with a generative model.</p> <ul style="list-style-type: none">● Workflow:<ul style="list-style-type: none">○ Convert documents into embeddings (vector representation).○ Store them in a vector database.○ On a query, retrieve top-k relevant chunks.○ Pass retrieved context to an LLM to generate an accurate, grounded answer.● Advantages:<ul style="list-style-type: none">○ Uses latest data without re-training model.

	<ul style="list-style-type: none"> ○ Reduces hallucination by grounding responses in retrieved knowledge.
Code	<pre> import os import numpy as np import google.generativeai as genai from typing import List, Dict, Any, Optional import requests from bs4 import BeautifulSoup import PyPDF2 import io from sentence_transformers import SentenceTransformer import json import pickle from dataclasses import dataclass from datetime import datetime import logging import faiss # Set up logging logging.basicConfig(level=logging.INFO) logger = logging.getLogger(__name__) @dataclass class Document: """Represents a document in the knowledge base""" content: str metadata: Dict[str, Any] id: str embedding: Optional[List[float]] = None class FAISSVectorStore: """FAISS-based vector storage for the RAG system""" def __init__(self, dimension: int = 384, index_path: str = "./faiss_index"): self.dimension = dimension self.index_path = index_path </pre>

```

        self.index = faiss.IndexFlatIP(dimension) # Inner product
similarity

        self.documents = {} # Store document content and metadata
        self.id_to_index = {} # Map document IDs to FAISS indices
        self.next_index = 0

        # Try to load existing index
        self.load_index()

        def add_documents(self, documents: List[str], embeddings:
List[List[float]],
                        metadatas: List[Dict], ids: List[str]):
            """Add documents to the FAISS index"""
            embeddings_array = np.array(embeddings, dtype=np.float32)

            # Normalize embeddings for cosine similarity
            faiss.normalize_L2(embeddings_array)

            # Add to FAISS index
            start_index = self.next_index
            self.index.add(embeddings_array)

            # Store documents and metadata
            for i, (doc, metadata, doc_id) in enumerate(zip(documents,
metadatas, ids)):
                current_index = start_index + i
                self.documents[current_index] = {
                    'content': doc,
                    'metadata': metadata,
                    'id': doc_id
                }
                self.id_to_index[doc_id] = current_index

            self.next_index += len(documents)
            self.save_index()

        def search(self, query_embedding: List[float], k: int = 5) ->
List[Dict]:
            """Search for similar documents"""

```

```

        if self.index.ntotal == 0:
            return []

        query_array = np.array([query_embedding], dtype=np.float32)
        faiss.normalize_L2(query_array)

        # Search
        scores, indices = self.index.search(query_array, min(k,
self.index.ntotal))

        results = []
        for score, idx in zip(scores[0], indices[0]):
            if idx >= 0 and idx in self.documents:
                result = self.documents[idx].copy()
                result['score'] = float(score)
                results.append(result)

        return results

def count(self) -> int:
    """Get total number of documents"""
    return self.index.ntotal

def save_index(self):
    """Save FAISS index and metadata to disk"""
    try:
        os.makedirs(os.path.dirname(self.index_path),
exist_ok=True)

        # Save FAISS index
        faiss.write_index(self.index, f"{self.index_path}.faiss")

        # Save metadata
        metadata = {
            'documents': self.documents,
            'id_to_index': self.id_to_index,
            'next_index': self.next_index,
            'dimension': self.dimension
        }

```

```

        with open(f"{self.index_path}_metadata.pkl", 'wb') as f:
            pickle.dump(metadata, f)

    except Exception as e:
        logger.error(f"Error saving index: {e}")

    def load_index(self):
        """Load FAISS index and metadata from disk"""
        try:
            index_file = f"{self.index_path}.faiss"
            metadata_file = f"{self.index_path}_metadata.pkl"

            if os.path.exists(index_file) and os.path.exists(metadata_file):
                # Load FAISS index
                self.index = faiss.read_index(index_file)

                # Load metadata
                with open(metadata_file, 'rb') as f:
                    metadata = pickle.load(f)
                    self.documents = metadata.get('documents', {})
                    self.id_to_index = metadata.get('id_to_index', {})
                    self.next_index = metadata.get('next_index', 0)

                logger.info(f"Loaded existing index with {self.index.ntotal} documents")
            else:
                logger.info("No existing index found, starting fresh")

        except Exception as e:
            logger.error(f"Error loading index: {e}")
            logger.info("Starting with fresh index")

class QuantumRAGAgent:
    def __init__(self, gemini_api_key: str = None):
        """
        Initialize the Quantum Computing RAG Agent

```

```

    Args:
        gemini_api_key: Gemini API key (will use .env if not
provided)
    """
    # Load API key
    if gemini_api_key:
        self.api_key = gemini_api_key
    else:
        self.api_key = secret_value_0

    if not self.api_key:
        raise ValueError("Gemini API key not found. Please set
GEMINI_API_KEY in .env or pass it directly.")

    # Configure Gemini
    genai.configure(api_key=self.api_key)
    self.model = genai.GenerativeModel('gemini-2.5-flash')

    # Initialize embedding model
    self.embedding_model = SentenceTransformer('all-MiniLM-L6-v2')

    # Initialize FAISS vector store
    self.vector_store = FAISSVectorStore(dimension=384) #
all-MiniLM-L6-v2 dimension

    logger.info("Quantum RAG Agent initialized successfully")

def extract_text_from_pdf(self, pdf_path: str) -> str:
    """Extract text from PDF file"""
    text = ""
    try:
        with open(pdf_path, 'rb') as file:
            pdf_reader = PyPDF2.PdfReader(file)
            for page in pdf_reader.pages:
                text += page.extract_text() + "\n"
    except Exception as e:
        logger.error(f"Error extracting text from PDF: {e}")
    return text

```

```

def extract_text_from_url(self, url: str) -> str:
    """Extract text from web URL"""
    try:
        response = requests.get(url, timeout=10)
        response.raise_for_status()
        soup = BeautifulSoup(response.content, 'html.parser')

        # Remove script and style elements
        for script in soup(["script", "style"]):
            script.decompose()

        # Get text and clean it
        text = soup.get_text()
        lines = (line.strip() for line in text.splitlines())
        chunks = (phrase.strip() for line in lines for phrase in
line.split("  "))
        text = ' '.join(chunk for chunk in chunks if chunk)

        return text
    except Exception as e:
        logger.error(f"Error extracting text from URL {url}: {e}")
        return ""

def chunk_text(self, text: str, chunk_size: int = 1000, overlap:
int = 200) -> List[str]:
    """Split text into overlapping chunks"""
    words = text.split()
    chunks = []

    for i in range(0, len(words), chunk_size - overlap):
        chunk = ' '.join(words[i:i + chunk_size])
        chunks.append(chunk)

    if i + chunk_size >= len(words):
        break

    return chunks

```

```

    def add_document(self, content: str, metadata: Dict[str, Any],
doc_id: str = None) -> str:
        """Add a document to the knowledge base"""
        if not doc_id:
            doc_id = f"doc_{datetime.now().timestamp()}"

        # Chunk the document
        chunks = self.chunk_text(content)

        # Prepare data for FAISS
        documents = []
        embeddings = []
        metadatas = []
        ids = []

        # Generate embeddings and prepare data
        for i, chunk in enumerate(chunks):
            chunk_id = f"{doc_id}_chunk_{i}"
            embedding = self.embedding_model.encode(chunk).tolist()

            documents.append(chunk)
            embeddings.append(embedding)
            metadatas.append(**metadata, "chunk_id": i,
"parent_doc_id": doc_id)
            ids.append(chunk_id)

        # Add to FAISS vector store
        self.vector_store.add_documents(documents, embeddings,
metadatas, ids)

        logger.info(f"Added document {doc_id} with {len(chunks)}
chunks")
        return doc_id

    def add_pdf(self, pdf_path: str, metadata: Dict[str, Any] = None)
-> str:
        """Add a PDF document to the knowledge base"""
        if not metadata:
            metadata = {"source": pdf_path, "type": "pdf"}

```



```

        text = self.extract_text_from_pdf(pdf_path)
        if not text.strip():
            logger.warning(f"No text extracted from PDF: {pdf_path}")
            return None

        return self.add_document(text, metadata)

    def add_url(self, url: str, metadata: Dict[str, Any] = None) ->
str:
        """Add content from a URL to the knowledge base"""
        if not metadata:
            metadata = {"source": url, "type": "web"}

        text = self.extract_text_from_url(url)
        if not text.strip():
            logger.warning(f"No text extracted from URL: {url}")
            return None

        return self.add_document(text, metadata)

    def add_text(self, text: str, metadata: Dict[str, Any] = None) ->
str:
        """Add raw text to the knowledge base"""
        if not metadata:
            metadata = {"type": "text", "added_at":
datetime.now().isoformat()}

        return self.add_document(text, metadata)

    def initialize_quantum_knowledge_base(self):
        """Initialize with basic quantum computing knowledge"""
        quantum_topics = [
            {
                "content": """
                    Quantum Computing Fundamentals:
                    Quantum computing is a revolutionary computational
                    paradigm that leverages quantum mechanical phenomena

```

such as superposition, entanglement, and interference to process information. Unlike classical bits

that can only be in states 0 or 1, quantum bits (qubits) can exist in superposition states, allowing them to be in multiple states simultaneously.

Key principles:

1. Superposition: Qubits can be in a combination of $|0\rangle$ and $|1\rangle$ states

2. Entanglement: Quantum particles can be correlated in ways that classical physics cannot explain

3. Interference: Quantum amplitudes can interfere constructively or destructively

4. Measurement: Observing a quantum system collapses it to a definite state

```
""",
    "metadata": {"topic": "fundamentals", "type":
"theory"}
```

```
},
```

```
{
```

```
"content": ""
```

Quantum Gates and Circuits:

Quantum gates are the building blocks of quantum circuits, analogous to logic gates in classical computing.

Common quantum gates include:

1. Pauli Gates (X, Y, Z): Single-qubit rotations around different axes

2. Hadamard Gate (H): Creates superposition by rotating $|0\rangle$ to $(|0\rangle + |1\rangle)/\sqrt{2}$

3. CNOT Gate: Two-qubit gate that creates entanglement

4. Phase Gates (S, T): Add phase to $|1\rangle$ state

5. Rotation Gates (RX, RY, RZ): Arbitrary rotations around X, Y, Z axes

Quantum circuits are represented as sequences of quantum gates applied to qubits,

read from left to right in time order.

```
""",
```

```

        "metadata": {"topic": "gates_circuits", "type":
"theory"}
    },
    {
        "content": ""
        Quantum Algorithms:
            Several quantum algorithms demonstrate quantum
advantage over classical algorithms:

            1. Shor's Algorithm: Efficiently factors large
integers, threatening RSA cryptography
            2. Grover's Algorithm: Searches unsorted databases
with quadratic speedup
            3. Quantum Fourier Transform: Essential subroutine for
many quantum algorithms
            4. Variational Quantum Eigensolver (VQE): Hybrid
quantum-classical algorithm for finding ground states
            5. Quantum Approximate Optimization Algorithm (QAOA):
For solving combinatorial optimization problems
            6. Deutsch-Jozsa Algorithm: Determines if a function
is constant or balanced

            These algorithms exploit quantum phenomena to achieve
computational advantages
            in specific problem domains.
        "",
        "metadata": {"topic": "algorithms", "type": "theory"}
    },
    {
        "content": ""
        Quantum Hardware Platforms:
            Different physical systems are used to implement
quantum computers:

            1. Superconducting Qubits: Used by IBM, Google,
Rigetti
            - Josephson junctions in superconducting circuits
            - Fast gate operations, but short coherence times

```

```

        2. Trapped Ion Systems: Used by IonQ, Honeywell
            - Individual ions trapped by electromagnetic fields
            - High fidelity, long coherence times, but slower
gates

        3. Photonic Systems: Used by Xanadu, PsiQuantum
            - Photons as qubits, room temperature operation
            - Natural for quantum communication

        4. Neutral Atoms: Used by QuEra, Pasqal
            - Atoms trapped by optical tweezers
            - Highly scalable architectures

        5. Silicon Spin Qubits: Emerging technology
            - Compatible with semiconductor manufacturing
        """
        "metadata": {"topic": "hardware", "type":
"technology"}
    }
]

for topic in quantum_topics:
    self.add_text(topic["content"], topic["metadata"])

logger.info("Initialized quantum computing knowledge base")

def retrieve_relevant_docs(self, query: str, n_results: int = 5)
-> List[Dict[str, Any]]:
    """Retrieve relevant documents for a query"""
    # Generate query embedding
    query_embedding = self.embedding_model.encode(query).tolist()

    # Search in FAISS vector store
    results = self.vector_store.search(query_embedding, n_results)

    return results

    def generate_response(self, query: str, context_docs:
List[Dict[str, Any]]) -> str:

```

```

        """Generate response using Gemini with retrieved context"""
        # Prepare context from retrieved documents
        context = "\n\n".join([
            f"[Source: {doc['metadata'].get('source', 'Knowledge
Base')}] \n{doc['content']}"
            for doc in context_docs
        ])

        # Shortened prompt to avoid token limits
        prompt = f"""Based on the quantum computing context provided,
answer this question: {query}

Context:
{context}

Provide a clear, accurate answer using the context above. Focus on the
key concepts and be concise."""

    try:
        # Add generation config to handle potential issues
        generation_config = genai.types.GenerationConfig(
            candidate_count=1,
            max_output_tokens=1000, # Limit output tokens
            temperature=0.3,
        )

        response = self.model.generate_content(
            prompt,
            generation_config=generation_config
        )

        # Check if response was blocked
        if response.candidates and
response.candidates[0].content.parts:
            return response.text
        else:
            logger.warning("Response was blocked or empty")
            return "I couldn't generate a response. The content
may have been filtered. Please try rephrasing your question."

```

```

        except Exception as e:
            logger.error(f"Error generating response: {e}")
            # More specific error handling
            if "quota" in str(e).lower() or "limit" in str(e).lower():
                return "API quota exceeded. Please check your Gemini
API usage limits."
            elif "safety" in str(e).lower() or "blocked" in
str(e).lower():
                return "Response was blocked by safety filters. Please
try rephrasing your question."
            else:
                return f"Error generating response: {str(e)}"

    def query(self, question: str, n_results: int = 5) -> Dict[str,
Any]:
        """Main query method - retrieve and generate response"""
        logger.info(f"Processing query: {question}")

        # Retrieve relevant documents
        relevant_docs = self.retrieve_relevant_docs(question,
n_results)

        if not relevant_docs:
            return {
                "query": question,
                "response": "No relevant documents found in the
knowledge base.",
                "sources": [],
                "retrieved_docs": 0
            }

        # Generate response
        response = self.generate_response(question, relevant_docs)

        return {
            "query": question,
            "response": response,
            "sources": [doc['metadata'] for doc in relevant_docs],

```

```

        "retrieved_docs": len(relevant_docs)
    }

def get_collection_stats(self) -> Dict[str, Any]:
    """Get statistics about the knowledge base"""
    count = self.vector_store.count()
    return {
        "total_chunks": count,
        "backend": "FAISS"
    }

# Example usage and testing
def main():
    """Example usage of the Quantum RAG Agent"""

    try:
        # Initialize the agent
        agent = QuantumRAGAgent()

        # Initialize with basic quantum knowledge
        agent.initialize_quantum_knowledge_base()

        # Example queries
        example_queries = [
            "What is quantum superposition?",
            "Explain Shor's algorithm",
            "What are quantum gates?",
            "How do superconducting qubits work?"
        ]

        print("Quantum Computing RAG Agent - Example Queries\n")
        print("=" * 50)

        for query in example_queries[:4]:
            print(f"\nQuery: {query}")
            print("-" * 30)

            result = agent.query(query)
            print(f"Response: {result['response']}")

```

	<pre>print(f"Sources used: {len(result['sources'])}") print(f"Retrieved chunks: {result['retrieved_docs']}") # Print collection stats stats = agent.get_collection_stats() print(f"\nKnowledge Base Stats:") print(f"Total chunks: {stats['total_chunks']}") print(f"Backend: {stats['backend']}") except Exception as e: print(f"Error in main: {e}") print("Make sure you have set your GEMINI_API_KEY in a .env file or environment variable") if __name__ == "__main__": main()</pre>
Output	<pre>Query: What is quantum superposition? ----- Batches: 100% ██████████ 1/1 [00:00<00:00, 39.73it/s] Response: Quantum superposition is a quantum mechanical phenomenon where quantum bits (qubits) can exist in multiple states simultaneously. Unlike classical bits that can only be in states 0 or 1, qubits in superposition can be in a combination of both the 0> and 1> states. The Hadamard Gate is an example of a quantum gate that can create superposition. Sources used: 4 Retrieved chunks: 4 Query: Explain Shor's algorithm ----- Batches: 100% ██████████ 1/1 [00:00<00:00, 66.84it/s] Response: Shor's Algorithm is a quantum algorithm that efficiently factors large integers. This capability demonstrates a quantum advantage over classical algorithms and poses a threat to RSA cryptography. Sources used: 4 Retrieved chunks: 4</pre>

	<div><div>Query: What are quantum gates? -----</div><div><div>Batches: 100%<div></div>1/1 [00:00<00:00, 51.59it/s]</div><div>Response: Quantum gates are the building blocks of quantum circuits, analogous to logic gates in classical computing. They are applied to qubits in a time-ordered sequence to perform operations.</div><div>Common quantum gates include:<ul style="list-style-type: none">* **Pauli Gates (X, Y, Z):** Perform single-qubit rotations around different axes.* **Hadamard Gate (H):** Creates superposition.* **CNOT Gate:** A two-qubit gate that creates entanglement.* **Phase Gates (S, T):** Add phase to the 1> state.* **Rotation Gates (RX, RY, RZ):** Perform arbitrary rotations around X, Y, Z axes.</div><div>Sources used: 4</div><div>Retrieved chunks: 4</div></div><div><div>Query: How do superconducting qubits work? -----</div><div><div>Batches: 100%<div></div>1/1 [00:00<00:00, 57.92it/s]</div><div>Response: Superconducting qubits work by utilizing Josephson junctions within superconducting circuits. These qubits are characterized by their fast gate operations, but they typically have short coherence times.</div><div>Sources used: 4</div><div>Retrieved chunks: 4</div></div></div></div>
Conclusion	<div>This experiment demonstrates how Retrieval Augmented Generation (RAG) can be used for specific tasks</div>